



CLOUDPANDITH ONLINE TRAININGS



Master Data Engineering with Expert-Led Live Training

◇ Azure Data Factory

Complete ETL/ELT

◇ Azure Databricks

Big Data & PySpark

◇ Microsoft Fabric

End-to-End Analytics

📞 ENROLL: **8904424822** | 🌐 www.cloudpandith.com

✓ Live Sessions ✓ Real Projects ✓ Expert Trainers ✓ Interview Prep ✓ Lifetime Access

- ◆ Live Classes
- ◆ Real-Time Projects
- ◆ Interview Preparation
- ◆ Lifetime Access

Index

1. Introduction to Schema Changes in Delta Lake
2. Schema Enforcement
3. Schema Evolution
4. Schema Enforcement vs Schema Evolution
5. Time Travel
6. Vacuum
7. Parquet File Cleanup vs Log Cleanup
8. Retention Summary
9. Complete Delta Lake Flow
10. Interview Questions and Answers

Delta Lake Features: Deep Clone vs Shallow Clone

Introduction

Delta Lake in Databricks provides advanced table management capabilities that help data engineers efficiently handle large-scale data workloads.

Two important Delta Lake features are:

- Deep Clone
- Shallow Clone

These features are used to create copies of Delta tables for:

- Backup and recovery
- Testing environments
- Data migration
- Sandbox experimentation
- Reporting and analytics
- CI/CD pipelines

Understanding the differences between Deep Clone and Shallow Clone is very important for designing scalable and cost-efficient data platforms.

What is Deep Clone?

A Deep Clone creates a fully independent copy of a Delta table.

It copies:

- Metadata
- Schema
- Transaction logs
- Parquet data files

After cloning, both tables become completely separate entities. Any modifications made to one table will not affect the other table.

Characteristics of Deep Clone

1. Complete Data Duplication

Deep Clone physically copies:

- All Parquet files
 - Delta transaction logs
 - Partition information
 - Metadata
-

2. Fully Independent Tables

After cloning:

- Source table and cloned table are completely isolated
 - Updates, deletes, and merges are independent
-

3. Reliable for Long-Term Usage

Deep Clone is ideal when:

- Data consistency is critical
 - Long-term backup is required
 - Independent environments are needed
-

4. Higher Storage Consumption

Since actual data files are copied:

- More cloud storage is required
 - Clone operation takes more time
-

Deep Clone Syntax

```
CREATE TABLE target_table  
DEEP CLONE source_table;
```

Deep Clone Example

```
CREATE TABLE sales_backup  
DEEP CLONE sales_data;
```

Explanation

- sales_data → Original table
 - sales_backup → Fully independent cloned table
-

Real-Time Use Cases of Deep Clone

1. Production Backup

Used for:

- Disaster recovery
 - Backup retention
 - Business continuity
-

2. Environment Migration

Useful for:

- Dev to QA migration
 - QA to Production migration
 - Cross-workspace migration
-

3. Historical Snapshot Preservation

Organizations use Deep Clone for:

- Auditing
 - Compliance reporting
 - Regulatory requirements
-

4. Safe Experimentation

Data scientists and engineers can:

- Test transformations
- Run experiments
- Validate ML models

without impacting production data.

What is Shallow Clone?

A Shallow Clone creates a metadata-only copy of a Delta table.

Instead of copying actual Parquet files, it references the same underlying data files used by the source table.

This makes Shallow Clone:

- Very fast
 - Storage efficient
 - Cost effective
-

Characteristics of Shallow Clone

1. Metadata Copy Only

Shallow Clone copies:

- Schema
 - Metadata
 - Partition information
 - Delta transaction logs
-

2. Shared Data Files

The source table and shallow cloned table reference the same Parquet files.

Data is NOT physically duplicated.

3. Faster Performance

Since no data copying occurs:

- Clone creation is extremely fast
 - Minimal compute resources are required
-

4. Lower Storage Cost

Because Parquet files are shared:

- Additional storage usage is very low
 - Cloud storage costs are reduced
-

Important Enhancement in Recent Delta Lake Versions

Earlier Delta Lake versions had an important limitation:

- VACUUM on source tables could remove Parquet files required by shallow clones
- This could break the shallow cloned table

However, in recent Delta Lake versions, this behavior has improved significantly.

Updated VACUUM Behavior

Modern Delta Lake versions maintain internal file reference tracking for shallow clones.

Now:

- VACUUM checks whether files are referenced by shallow clone tables
- Referenced Parquet files are protected
- Files required by shallow clones are NOT deleted

Example:

VACUUM source_table;

Even after VACUUM:

- Shared files used by shallow clone remain safe
- Clone continues to function properly

This improvement makes shallow clone much more reliable and production-friendly.

Shallow Clone Syntax

```
CREATE TABLE target_table  
SHALLOW CLONE source_table;
```

Shallow Clone Example

```
CREATE TABLE sales_analysis  
SHALLOW CLONE sales_data;
```

Explanation

- Only metadata is copied
 - Data files remain shared with source table
-

Real-Time Use Cases of Shallow Clone

1. Fast Development Environments

Used for:

- Sandbox environments
 - Developer testing
 - Feature validation
-

2. Temporary Analytics

Useful for:

- Ad-hoc analysis
 - Reporting
 - Data exploration
-

3. CI/CD Pipelines

Teams use shallow clone for:

- Automated testing
- Temporary validation environments
- Deployment verification

4. Schema Testing

Engineers can:

- Test schema modifications
- Validate partitions
- Perform metadata experimentation

without duplicating large datasets.

5. Cost Optimization

Shallow Clone helps organizations:

- Reduce storage costs
 - Avoid unnecessary duplication
 - Provision environments quickly
-

Deep Clone vs Shallow Clone Comparison

Feature	Shallow Clone	Deep Clone
Data Copy	No	Yes
Metadata Copy	Yes	Yes
Parquet Files	Shared	Separate
Storage Usage	Low	High
Performance	Very Fast	Slower
Table Independence	Partial	Complete
Backup Capability	Limited	Excellent
Environment Provisioning	Excellent	Good
Cost Efficiency	High	Lower
VACUUM Protection	Supported in recent versions	Fully independent

Architecture-Level Understanding

Deep Clone Architecture

- Copies complete storage layer
- Creates independent Delta logs
- Creates separate Parquet files

Result:

- Full isolation between tables
-

Shallow Clone Architecture

- Creates new metadata layer
- References existing Parquet files
- Shares physical storage

Result:

- Faster creation
 - Lower storage usage
-

Best Practices

Use Deep Clone When

- Long-term backup is required
 - Independent data copies are needed
 - Compliance and auditing exist
 - Production isolation is important
 - Cross-region migration is required
-

Use Shallow Clone When

- Fast cloning is needed
- Storage optimization matters

- Temporary environments are required
 - CI/CD testing pipelines are used
 - Quick reporting environments are needed
-

Important Considerations

Deep Clone

Advantages

- Complete isolation
- Reliable backup
- Safer for production workloads

Limitations

- Higher storage cost
 - Longer execution time
-

Shallow Clone

Advantages

- Very fast
- Low storage usage
- Cost efficient

Limitations

- Shared storage dependency
 - Not fully independent like deep clone
-

Conclusion

Deep Clone and Shallow Clone are powerful Delta Lake features that solve different business requirements.

Deep Clone

Best suited for:

- Backup
- Disaster recovery
- Independent environments
- Long-term retention

Shallow Clone

Best suited for:

- Fast provisioning
- Temporary environments
- Analytics
- CI/CD pipelines
- Cost optimization

Recent improvements in Delta Lake have made Shallow Clone significantly more reliable by protecting referenced Parquet files during VACUUM operations.

Choosing the right cloning strategy depends on:

- Performance requirements
- Storage cost
- Isolation needs
- Backup requirements
- Operational complexity